

## hw3

April 29, 2016

```
In [1]: %pylab inline
        %matplotlib inline
```

Populating the interactive namespace from numpy and matplotlib

```
In [165]: import numpy as np
import pandas as pd
import scipy.stats as ss
import matplotlib.pyplot as plt
from scipy.stats import norm

import math
from ctypes import *
import sys
import string
```

```
In [166]: review_df = pd.read_csv('2015 CHR Analytic Data.csv')
len(review_df)
```

```
Out[166]: 3191
```

```
In [167]: #Step 1
```

```
df = review_df # Store in dummy variable
df['2011 population estimate Value'] = df["2011 population estimate Value"].str.replace(',','')
df['2011 population estimate Value'] = df['2011 population estimate Value'].astype(float)
df = df.drop(df[df['2011 population estimate Value'] < 30000].index) # Drop all < 30000
#print len(df)
```

```
#Step 2
```

```
for column_name, column in df.transpose().iterrows():
    if(column_name!='COUNTYCODE' and column_name.find("Value",(len(column_name)-5))== -1):
        df = df.drop(column_name, 1)
#print len(df)
```

```
#Step3
```

```
df= df.dropna()
```

```
#Delete all total rows
```


```
df.dropna() #Drop all rows having any Nan
df = df.drop(df[df['COUNTYCODE'] == 0].index)
print "1. TOTAL NUMBER OF COUNTIES: ",len(df)
```

1. TOTAL NUMBER OF COUNTIES: 823

```
In [205]: df['log_paamv'] = df['Premature age-adjusted mortality Value']
```

```
df['log_paamv']=np.log(df['log_paamv'])
df['log_paamv'].hist(bins=20)
print "2. log_paamv HISTOGRAM: 2histogram.png"
```

2. log\_paamv HISTOGRAM: 2histogram.png



```
In [182]: from sklearn import datasets, linear_model
```

```
#Shuffle Indexes
df = df.reindex(np.random.permutation(df.index))
#print df.head()
```

```
REMOVED_COLUMNS = ['COUNTYCODE', 'log_paamv', 'Premature age-adjusted mortality Value', 'Prema
lr_df = df.copy()
```

```
for col_name in REMOVED_COLUMNS:
    #print col_name
    lr_df = lr_df.drop(col_name, 1)
lr_df.head()
```

```
for c in lr_df.columns:
    if not (lr_df[c].dtype == np.float64 or lr_df[c].dtype == np.int64):
        lr_df[c] = lr_df[c].apply(lambda val: float(string.replace(str(val), ',', ''))) ##chang
        lr_df[c] = lr_df[c].astype('float')
```

```

#Standardize

regr = linear_model.LinearRegression()

gMAE = 0.0
for j in range(0, 10):

    to_be_skipped_from = int( (j*len(lr_df))/10 )
    to_be_skipped = int(len(lr_df)/10)+1

    train_x = lr_df[0:to_be_skipped_from].copy()
    train_y = df['log_paamv'][0:to_be_skipped_from].copy()

    train_x = train_x.append(lr_df[to_be_skipped_from+to_be_skipped:len(lr_df)+1].copy())
    train_y = train_y.append(df['log_paamv'][to_be_skipped_from+to_be_skipped:len(lr_df)+1].copy())

    test_x = lr_df[to_be_skipped_from:to_be_skipped_from+to_be_skipped].copy()
    test_y = df['log_paamv'][to_be_skipped_from:to_be_skipped_from+to_be_skipped].copy()

    for c in train_x.columns:
        train_x[c] = (train_x[c] - np.mean(train_x[c]))/np.std(train_x[c])

    train_y = (train_y - np.mean(train_y))/np.std(train_y)

    for c in test_x.columns:
        test_x[c] = (test_x[c] - np.mean(test_x[c]))/np.std(test_x[c])

    test_y = (test_y - np.mean(test_y))/np.std(test_y)

    #print 'train_x',0,to_be_skipped_from
    #print 'train_x append',to_be_skipped_from+to_be_skipped,len(lr_df)+1
    #print 'test_x',to_be_skipped_from,to_be_skipped_from+to_be_skipped

    regr.fit(train_x, train_y)

    # The coefficients
    #print("Residual sum of squares:",np.mean((regr.predict(test_x) - test_y) ** 2))
    gMAE += np.mean((regr.predict(test_x) - test_y) ** 2)

gMAE = gMAE/10

print "3. Non-regularized Linear Regression MSE: ", gMAE

```

3. Non-regularized Linear Regression MSE: 0.124838899866

In [203]: *#Part 4*

```

from sklearn.decomposition import PCA

#Standardize dataset
for c in lr_df.columns:
    lr_df[c] = (lr_df[c] - np.mean(lr_df[c]))/np.std(lr_df[c])

pca = PCA(n_components=3)
pca.fit(lr_df)

```

```

print "4. Percentage variance explained of first three components:", pca.explained_variance_r

4. Percentage variance explained of first three components: [ 0.28090638  0.13937027  0.0792126 ]

In [199]: #Part 5 a)
#PCA
from sklearn import datasets, linear_model

#Shuffle Indexes
df = df.reindex(np.random.permutation(df.index))
#print df.head()

lr_df = df.copy()
#Standardize

for col_name in REMOVED_COLUMNS:
    #print col_name
    lr_df = lr_df.drop(col_name, 1)
lr_df.head()

for c in lr_df.columns:
    if not (lr_df[c].dtype == np.float64 or lr_df[c].dtype == np.int64):
        lr_df[c] = lr_df[c].apply(lambda val: float(string.replace(str(val),',',''))) ##chang
        lr_df[c] = lr_df[c].astype('float')

regr = linear_model.LinearRegression()
gMAE = 0.0

for j in range(0, 10):

    to_be_skipped_from = int( (j*len(lr_df))/10 )
    to_be_skipped = int(2*len(lr_df)/10)+1

    if j < 9:
        train_x = lr_df[0:to_be_skipped_from].copy()
        train_y = df['log_paamv'][0:to_be_skipped_from].copy()

        train_x = train_x.append(lr_df[to_be_skipped_from+to_be_skipped:len(lr_df)+1].copy())
        train_y = train_y.append(df['log_paamv'][to_be_skipped_from+to_be_skipped:len(lr_df)+1].copy())

        test_x = lr_df[to_be_skipped_from:to_be_skipped_from+to_be_skipped/2].copy()
        test_y = df['log_paamv'][to_be_skipped_from:to_be_skipped_from+to_be_skipped/2].copy()

        dev_x = lr_df[to_be_skipped_from+to_be_skipped/2:to_be_skipped_from+to_be_skipped].copy()
        dev_y = df['log_paamv'][to_be_skipped_from+to_be_skipped/2:to_be_skipped_from+to_be_skipped].copy()

    else:
        to_be_skipped = int(len(lr_df)/10)+1
        train_x = lr_df[to_be_skipped+1:len(lr_df)-to_be_skipped].copy()
        train_y = df['log_paamv'][to_be_skipped+1:len(lr_df)-to_be_skipped].copy()

        test_x = lr_df[len(lr_df)-to_be_skipped+1:len(lr_df)].copy()

```

```

test_y = df['log_paamv'][len(lr_df)-to_be_skipped+1:len(lr_df)].copy()

dev_x = lr_df[0:to_be_skipped].copy()
dev_y = df['log_paamv'][0:to_be_skipped].copy()

for c in train_x.columns:
    train_x[c] = (train_x[c] - np.mean(train_x[c]))/np.std(train_x[c])

train_y = (train_y - np.mean(train_y))/np.std(train_y)

for c in test_x.columns:
    test_x[c] = (test_x[c] - np.mean(test_x[c]))/np.std(test_x[c])

test_y = (test_y - np.mean(test_y))/np.std(test_y)

for c in dev_x.columns:
    dev_x[c] = (dev_x[c] - np.mean(dev_x[c]))/np.std(dev_x[c])

dev_y = (dev_y - np.mean(dev_y))/np.std(dev_y)

#print 'train_x',0,to_be_skipped_from
#print 'train_x append',to_be_skipped_from+to_be_skipped,len(lr_df)+1
#print 'test_x',to_be_skipped_from,to_be_skipped_from+to_be_skipped
min_error = 10000;
min_comp = 1
for k in range(3,44):
    pca = PCA(n_components=k)
    pca.fit(train_x)
    Z = pca.transform(train_x)
    Z = pd.DataFrame(Z).copy()
    dev_min = pca.transform(dev_x)
    dev_min = pd.DataFrame(dev_min).copy()
    #print 'Z: ', len(Z), ' train_y', len(train_y), ' train_x', len(train_x)
    #print 'dev_x', len(dev_x), ' dev_y', len(dev_y)
    regr.fit(Z, train_y)
    error = np.mean((regr.predict(dev_min) - dev_y) ** 2)
    if error < min_error:
        min_comp = k
        min_error = error
#print 'iter ',j
#print 'min components', min_comp
pca = PCA(n_components=min_comp)
pca.fit(train_x)
Z = pca.transform(train_x)
Z = pd.DataFrame(Z).copy()
test_x = pca.transform(test_x)
test_x = pd.DataFrame(test_x).copy()
regr.fit(Z, train_y)
gMAE += np.mean((regr.predict(test_x) - test_y) ** 2)

gMAE = gMAE/10

print "5. a) principal components regression mse: ", gMAE

```

5. a) principal components regression mse: 0.126886957947

In [201]: *#Part 5 b)*

*#L2 Regularization*

```
from sklearn.linear_model import Ridge, Lasso
```

*#Shuffle Indexes*

```
df = df.reindex(np.random.permutation(df.index))
```

*#print df.head()*

```
lr_df = df.copy()
```

*#Standardize*

```
for col_name in REMOVED_COLUMNS:
```

*#print col\_name*

```
    lr_df = lr_df.drop(col_name, 1)
```

```
lr_df.head()
```

```
for c in lr_df.columns:
```

```
    if not (lr_df[c].dtype == np.float64 or lr_df[c].dtype == np.int64):
```

```
        lr_df[c] = lr_df[c].apply(lambda val: float(string.replace(str(val),',',''))) ##change
```

```
        lr_df[c] = lr_df[c].astype('float')
```

```
gMAE = 0.0
```

```
for j in range(0, 10):
```

```
    to_be_skipped_from = int( (j*len(lr_df))/10 )
```

```
    to_be_skipped = int(2*len(lr_df)/10)+1
```

```
    if j < 9:
```

```
        train_x = lr_df[0:to_be_skipped_from].copy()
```

```
        train_y = df['log_paamv'][0:to_be_skipped_from].copy()
```

```
        train_x = train_x.append(lr_df[to_be_skipped_from+to_be_skipped:len(lr_df)+1].copy())
```

```
        train_y = train_y.append(df['log_paamv'][to_be_skipped_from+to_be_skipped:len(lr_df)+1].copy())
```

```
        test_x = lr_df[to_be_skipped_from:to_be_skipped_from+to_be_skipped/2].copy()
```

```
        test_y = df['log_paamv'][to_be_skipped_from:to_be_skipped_from+to_be_skipped/2].copy()
```

```
        dev_x = lr_df[to_be_skipped_from+to_be_skipped/2:to_be_skipped_from+to_be_skipped].copy()
```

```
        dev_y = df['log_paamv'][to_be_skipped_from+to_be_skipped/2:to_be_skipped_from+to_be_skipped].copy()
```

```
    else:
```

```
        to_be_skipped = int(len(lr_df)/10)+1
```

```
        train_x = lr_df[to_be_skipped+1:len(lr_df)-to_be_skipped].copy()
```

```
        train_y = df['log_paamv'][to_be_skipped+1:len(lr_df)-to_be_skipped].copy()
```

```
        test_x = lr_df[len(lr_df)-to_be_skipped+1:len(lr_df)].copy()
```

```
        test_y = df['log_paamv'][len(lr_df)-to_be_skipped+1:len(lr_df)].copy()
```

```
        dev_x = lr_df[0:to_be_skipped].copy()
```

```

dev_y = df['log_paamv'][0:to_be_skipped].copy()

for c in train_x.columns:
    train_x[c] = (train_x[c] - np.mean(train_x[c]))/np.std(train_x[c])

train_y = (train_y - np.mean(train_y))/np.std(train_y)

for c in test_x.columns:
    test_x[c] = (test_x[c] - np.mean(test_x[c]))/np.std(test_x[c])

test_y = (test_y - np.mean(test_y))/np.std(test_y)

for c in dev_x.columns:
    dev_x[c] = (dev_x[c] - np.mean(dev_x[c]))/np.std(dev_x[c])

dev_y = (dev_y - np.mean(dev_y))/np.std(dev_y)

#print 'train_x',0,to_be_skipped_from
#print 'train_x append',to_be_skipped_from+to_be_skipped,len(lr_df)+1
#print 'test_x',to_be_skipped_from,to_be_skipped_from+to_be_skipped
min_error = 10000;
min_alpha = 1
alphas = [pow(10,-5),pow(10,-4),pow(10,-3),pow(10,-2),pow(10,-1),1,10,100,1000,10000,100000]
for alpha in alphas:
    regr = Ridge(alpha=alpha)
    regr.fit(train_x, train_y)
    error = np.mean((regr.predict(dev_x) - dev_y) ** 2)
    if error < min_error:
        min_alpha = alpha
        min_error = error
#print 'iter ',j
#print 'min alpha', min_alpha
#print 'min error', min_error
regr = Ridge(alpha=min_alpha)
regr.fit(train_x, train_y)
gMAE += np.mean((regr.predict(test_x) - test_y) ** 2)

gMAE = gMAE/10

print "5. b) L2 regularized mse: ", gMAE

```

5. b) L2 regularized mse: 0.12111448232

In [202]: #Part 5 c)

*#L1 Regularization*

```
from sklearn.linear_model import Ridge, Lasso
```

*#Shuffle Indexes*

```
df = df.reindex(np.random.permutation(df.index))
```

```
#print df.head()
```

```
lr_df = df.copy()
```

*#Standardize*

```

for col_name in REMOVED_COLUMNS:
    #print col_name
    lr_df = lr_df.drop(col_name, 1)
lr_df.head()

for c in lr_df.columns:
    if not (lr_df[c].dtype == np.float64 or lr_df[c].dtype == np.int64):
        lr_df[c] = lr_df[c].apply(lambda val: float(string.replace(str(val), ',', ''))) ##chang
        lr_df[c] = lr_df[c].astype('float')

gMAE = 0.0

for j in range(0, 10):

    to_be_skipped_from = int( (j*len(lr_df))/10 )
    to_be_skipped = int(2*len(lr_df)/10)+1

    if j < 9:
        train_x = lr_df[0:to_be_skipped_from].copy()
        train_y = df['log_paamv'][0:to_be_skipped_from].copy()

        train_x = train_x.append(lr_df[to_be_skipped_from+to_be_skipped:len(lr_df)+1].copy())
        train_y = train_y.append(df['log_paamv'][to_be_skipped_from+to_be_skipped:len(lr_df)+1].copy())

        test_x = lr_df[to_be_skipped_from:to_be_skipped_from+to_be_skipped/2].copy()
        test_y = df['log_paamv'][to_be_skipped_from:to_be_skipped_from+to_be_skipped/2].copy()

        dev_x = lr_df[to_be_skipped_from+to_be_skipped/2:to_be_skipped_from+to_be_skipped].copy()
        dev_y = df['log_paamv'][to_be_skipped_from+to_be_skipped/2:to_be_skipped_from+to_be_skipped].copy()

    else:
        to_be_skipped = int(len(lr_df)/10)+1
        train_x = lr_df[to_be_skipped+1:len(lr_df)-to_be_skipped].copy()
        train_y = df['log_paamv'][to_be_skipped+1:len(lr_df)-to_be_skipped].copy()

        test_x = lr_df[len(lr_df)-to_be_skipped+1:len(lr_df)].copy()
        test_y = df['log_paamv'][len(lr_df)-to_be_skipped+1:len(lr_df)].copy()

        dev_x = lr_df[0:to_be_skipped].copy()
        dev_y = df['log_paamv'][0:to_be_skipped].copy()

    for c in train_x.columns:
        train_x[c] = (train_x[c] - np.mean(train_x[c]))/np.std(train_x[c])

    train_y = (train_y - np.mean(train_y))/np.std(train_y)

    for c in test_x.columns:
        test_x[c] = (test_x[c] - np.mean(test_x[c]))/np.std(test_x[c])

    test_y = (test_y - np.mean(test_y))/np.std(test_y)

    for c in dev_x.columns:

```



```

dev_x[c] = (dev_x[c] - np.mean(dev_x[c]))/np.std(dev_x[c])

dev_y = (dev_y - np.mean(dev_y))/np.std(dev_y)

#print 'train_x',0,to_be_skipped_from
#print 'train_x append',to_be_skipped_from+to_be_skipped,len(lr_df)+1
#print 'test_x',to_be_skipped_from,to_be_skipped_from+to_be_skipped
min_error = 10000;
min_alpha = 1
alphas = [pow(10,-5),pow(10,-4),pow(10,-3),pow(10,-2),pow(10,-1),1,10,100,1000,10000,100000]
for alpha in alphas:
    regr = Lasso(alpha=alpha)
    regr.fit(train_x, train_y)
    error = np.mean((regr.predict(dev_x) - dev_y) ** 2)
    if error < min_error:
        min_alpha = alpha
        min_error = error
#print 'iter ',j
#print 'min alpha', min_alpha
#print 'min error', min_error
regr = Lasso(alpha=min_alpha)
regr.fit(train_x, train_y)
gMAE += np.mean((regr.predict(test_x) - test_y) ** 2)

gMAE = gMAE/10

print "5. c) L1 regularized mse: ", gMAE

```

5. b) L1 regularized mse: 0.122738420223

In [ ]: